

# END TERM PAPER 2024 SOLUTION DBMD

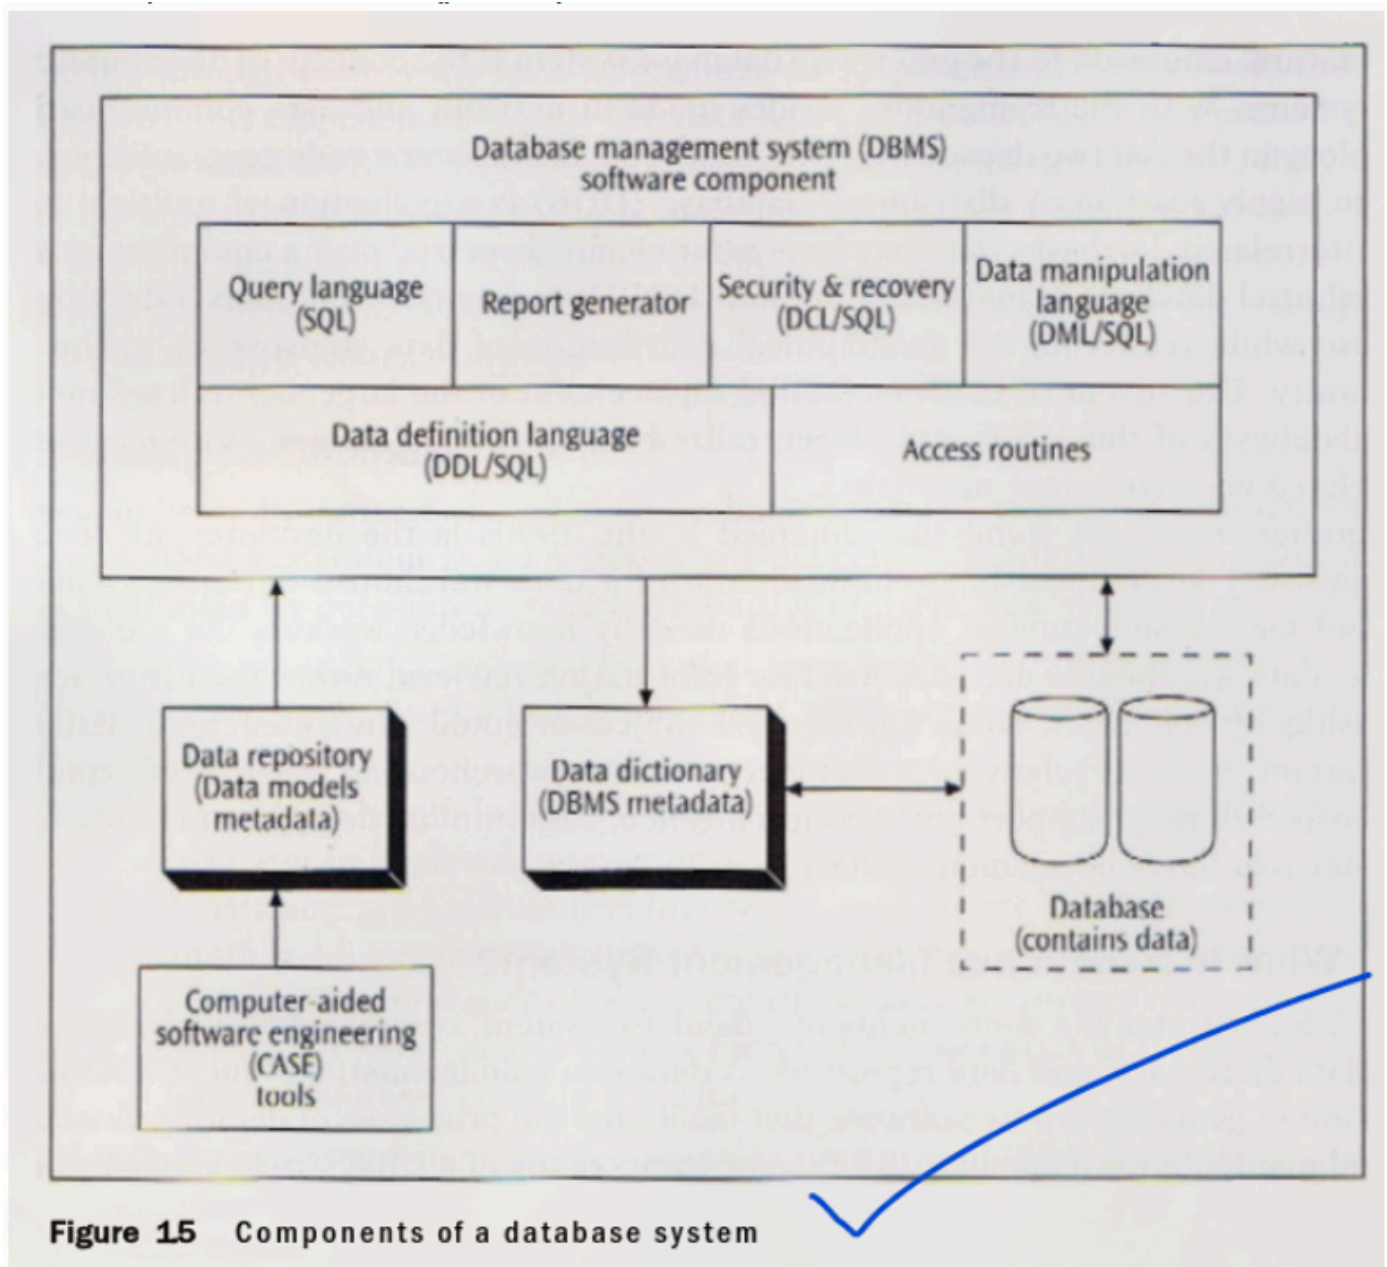
## (SIMPLIFIED)

### Question 1 (Compulsory)

(Attempt all parts - 3x5 = 15 Marks)

a) Describe the components of database systems. (3 Marks)

Answer:



A database system consists of integrated components working together to manage data:

1. **Hardware:** Physical devices like servers, storage disks, and network devices.
2. **Software:**

- **Database Management System (DBMS):** Core software to define, create, manipulate, control, and manage the database. Includes:
  - **Query Processor:** Interprets and executes queries (DDL interpreter, DML compiler, query evaluation engine).
  - **Storage Manager:** Manages physical storage (buffer, file, transaction management).
- **Application Programs & Utilities:** Software for data access, manipulation, reporting, backups (e.g., report generators, data import/export tools).

3. **Data:** Stored information, including operational data and metadata (data about data, in a data dictionary).

4. **Users:**

- **Database Administrators (DBAs):** Manage the overall system (security, backup, recovery, performance).
- **Database Designers:** Define database structure (schema).
- **Application Programmers:** Develop applications interacting with the database.
- **End Users:** Access data for queries, updates, and reports.

5. **Procedures:** Instructions and rules for database design and use (e.g., login, backup/recovery methods).

(Based on: DBMD UNIT - 1 and 2 Notes, p.3, Figure 1.5 and related text; general DBMS knowledge)

**b) Distinguish between primary key, candidate key and super key. (3 Marks)**

**Answer:**

Feature	Superkey	Candidate Key	Primary Key
<b>Definition</b>	Any attribute(s) uniquely identifying a tuple (row).	A minimal superkey; no attribute can be removed without losing uniqueness.	The candidate key chosen by the designer to uniquely identify tuples.
<b>Minimality</b>	Not necessarily minimal; may contain redundant attributes.	Must be minimal.	Must be minimal (as it's a candidate key).
<b>Uniqueness</b>	Guarantees uniqueness.	Guarantees uniqueness.	Guarantees uniqueness.
<b>Number</b>	Many per relation.	One or more per relation.	Only one per relation.
<b>Null Values</b>	Generally should not be null (PK enforces no nulls).	Attributes should ideally not be null.	Cannot contain null values (Entity Integrity Constraint).
<b>Example</b>	If {SID, CID} is a candidate key, then	In Employee table, {EmpID} and {SSN}	From {EmpID}, {SSN}, DBA might choose

Feature	Superkey	Candidate Key	Primary Key
	{SID, CID, SName} is a superkey.	could be candidate keys.	{EmpID} as primary key.

(Based on: DBMD UNIT - 1 and 2 Notes, p.8 "Unique Identifiers (Keys)", p.42 "Key Constraints")

### c) Explain the advanced data manipulation using SQL. (3 Marks)

#### Answer:

Advanced Data Manipulation Language (DML) in SQL extends beyond basic `INSERT`, `SELECT`, `UPDATE`, `DELETE` for single rows, involving:

#### 1. Complex Queries:

- **Joins:** Combining data from multiple tables (e.g., `INNER JOIN`, `LEFT JOIN`).
- **Subqueries (Nested Queries):** Queries within other SQL queries for complex filtering or calculations, including correlated subqueries.
- **Aggregate Functions with Grouping:** Using `COUNT()`, `SUM()`, etc., with `GROUP BY` for group calculations, and `HAVING` to filter groups.

#### 2. Set Operations:

Combining `SELECT` results using `UNION`, `UNION ALL`, `INTERSECT`, `EXCEPT` (or `MINUS`).

#### 3. Window Functions (Advanced):

Calculations across a set of table rows related to the current row (e.g., ranking, moving averages) without collapsing rows like `GROUP BY`.

#### 4. Common Table Expressions (CTEs):

Using `WITH` to define temporary, named result sets for readability and modularity in complex queries.

#### 5. Bulk Operations:

Efficiently inserting, updating, or deleting large data volumes (e.g., `INSERT INTO ... SELECT ...`).

These features enable sophisticated data retrieval, analysis, and modification for reporting, business intelligence, and complex applications.

(Based on: DBMD UNIT - 3 and 4 Notes, p.12, p.15; general SQL knowledge extending beyond basic DML definitions in Unit 1 notes)

### d) Explain the Data Control Language (DCL) commands. (3 Marks)

#### Answer:

Data Control Language (DCL) commands in SQL manage permissions and control access to database objects. Main DCL commands:

#### 1. GRANT:

- **Purpose:** Gives specific privileges (e.g., `SELECT`, `INSERT`) on database objects (tables, views) to users or roles.
- **Syntax (Simplified):** `GRANT privilege_list ON object_name TO user_list [WITH GRANT OPTION];`
- **WITH GRANT OPTION:** Allows grantee to grant received privileges to others.
- **Example:** `GRANT SELECT, INSERT ON Employees TO 'john_doe';`


## 2. REVOKE:

- **Purpose:** Removes previously granted or denied privileges from users or roles.
- **Syntax (Simplified):** `REVOKE [GRANT OPTION FOR] privilege_list ON object_name FROM user_list [CASCADE | RESTRICT];`
- **GRANT OPTION FOR:** Revokes only the grant ability, not the privilege itself.
- **CASCADE:** Revokes privilege from user and from others to whom this user granted it.
- **Example:** `REVOKE INSERT ON Employees FROM 'john_doe';`

DCL commands are vital for database security, ensuring users perform only authorized actions.

(Based on: DBMD UNIT - 1 and 2 Notes, p.4; DBMD UNIT - 3 and 4 Notes, p.32-34)

e) Find the minimum number of tables required to represent the given ER diagram in the relational model. (3 Marks)



**Answer:**

Mapping ER diagram to tables:

- Entity M (Strong):** Attributes M1 (PK), M2, M3.
  - Rule: Strong entity -> own table.
  - **Table 1: Table\_M (M1, M2, M3)**
- Entity P (Strong):** Attributes P1 (PK), P2.
  - Rule: Strong entity -> own table.
  - **Table 2: Table\_P (P1, P2)**
- Entity N (Weak):** Attributes N1 (Partial Key), N2. Identified by P via R2.
  - Rule: Weak entity -> own table. PK = PK of owner (P) + partial key (N1).

- **Table 3: Table\_N (P1\_FK, N1, N2)** (P1\_FK references P)

#### 4. Relationship R1 (M to P, 1:N):

- Rule (1:N): PK of "1" side (M1 from M) becomes FK in "N" side table (P).
- Table\_P updated: **Table\_P (P1, P2, M1\_FK)** (M1\_FK references M)
- No new table for 1:N.

#### 5. Relationship R2 (P to N, M:1, Identifying):

- Rule (Identifying): Already handled by including P's PK (P1\_FK) in Table\_N's composite PK.
- No additional table needed.

#### Consolidated Tables:

1. **Table\_M (M1, M2, M3)**
2. **Table\_P (P1, P2, M1\_FK)** (M1\_FK references Table\_M)
3. **Table\_N (P1\_FK, N1, N2)** (P1\_FK references Table\_P; {P1\_FK, N1} is PK)

Minimum number of tables required is **3**.

*(Based on: ER to Relational Mapping rules covered in DBMD UNIT - 1 and 2 Notes, p.43-44)*

---

#### UNIT-I

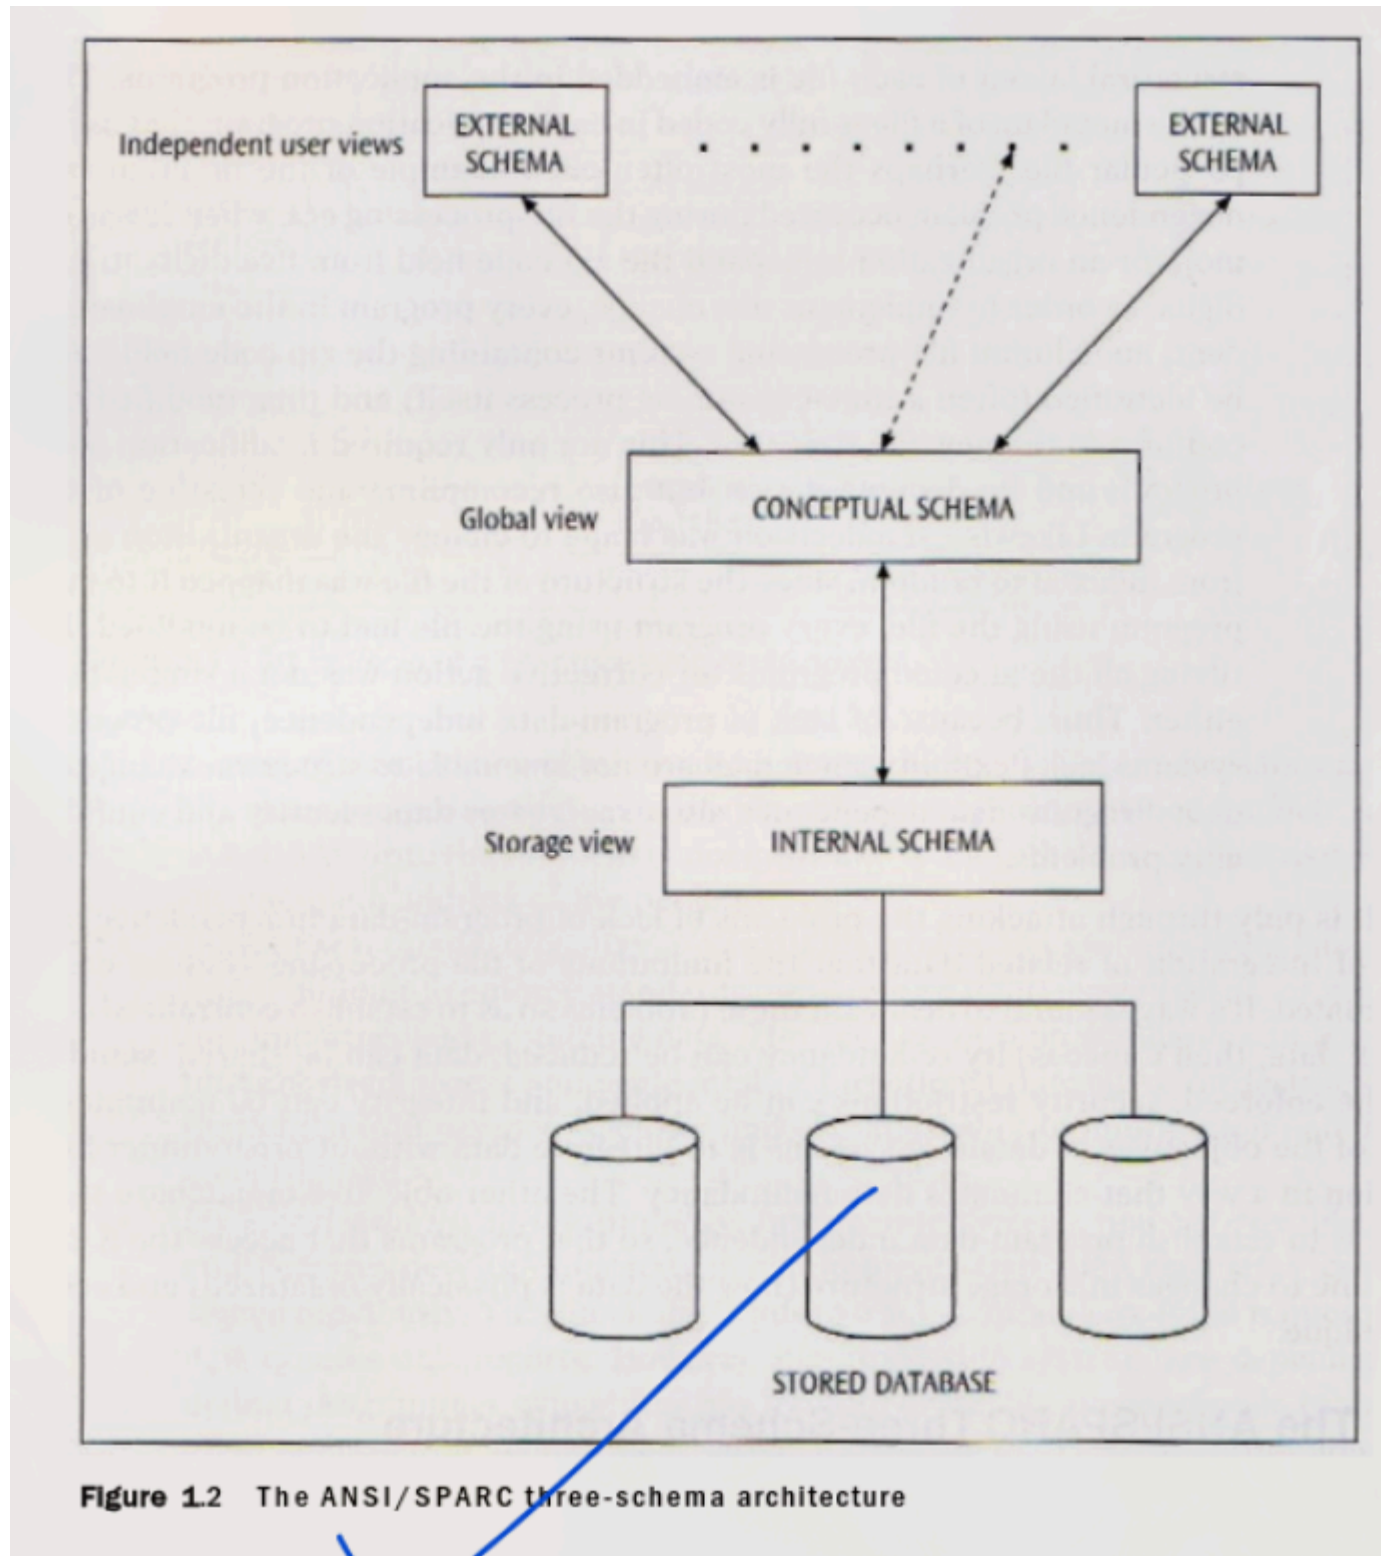
(Select one question from Q2 or Q3)

#### Question 2:

**a) Explain the database systems architecture with a suitable diagram. (8 Marks)**



Answer:



The ANSI/SPARC three-schema architecture promotes data independence by separating user views from physical storage. It has three levels:

**1. External Level (User Views / Subschemas):**

- Closest to users; describes the database part relevant to a specific user/group.
- Different users can have different views (e.g., HR sees salaries, project managers see skills).
- Defined by an external schema; provides security by hiding irrelevant/sensitive data.

**2. Conceptual Level (Global View / Conceptual Schema):**

- Community view of the entire database; describes its logical structure for all users.
- Defines entities, attributes, relationships, constraints, and semantic information.
- Technology-independent; hides physical storage details. Database designers work at this level.

### 3. Internal Level (Physical View / Internal Schema):

- Lowest level, closest to physical storage; describes how data is physically stored.
- Specifies data structures, file organizations (e.g., B+-trees), access paths (indexes), compression, encryption.
- Technology-dependent; deals with storage efficiency and access.

### Data Independence:

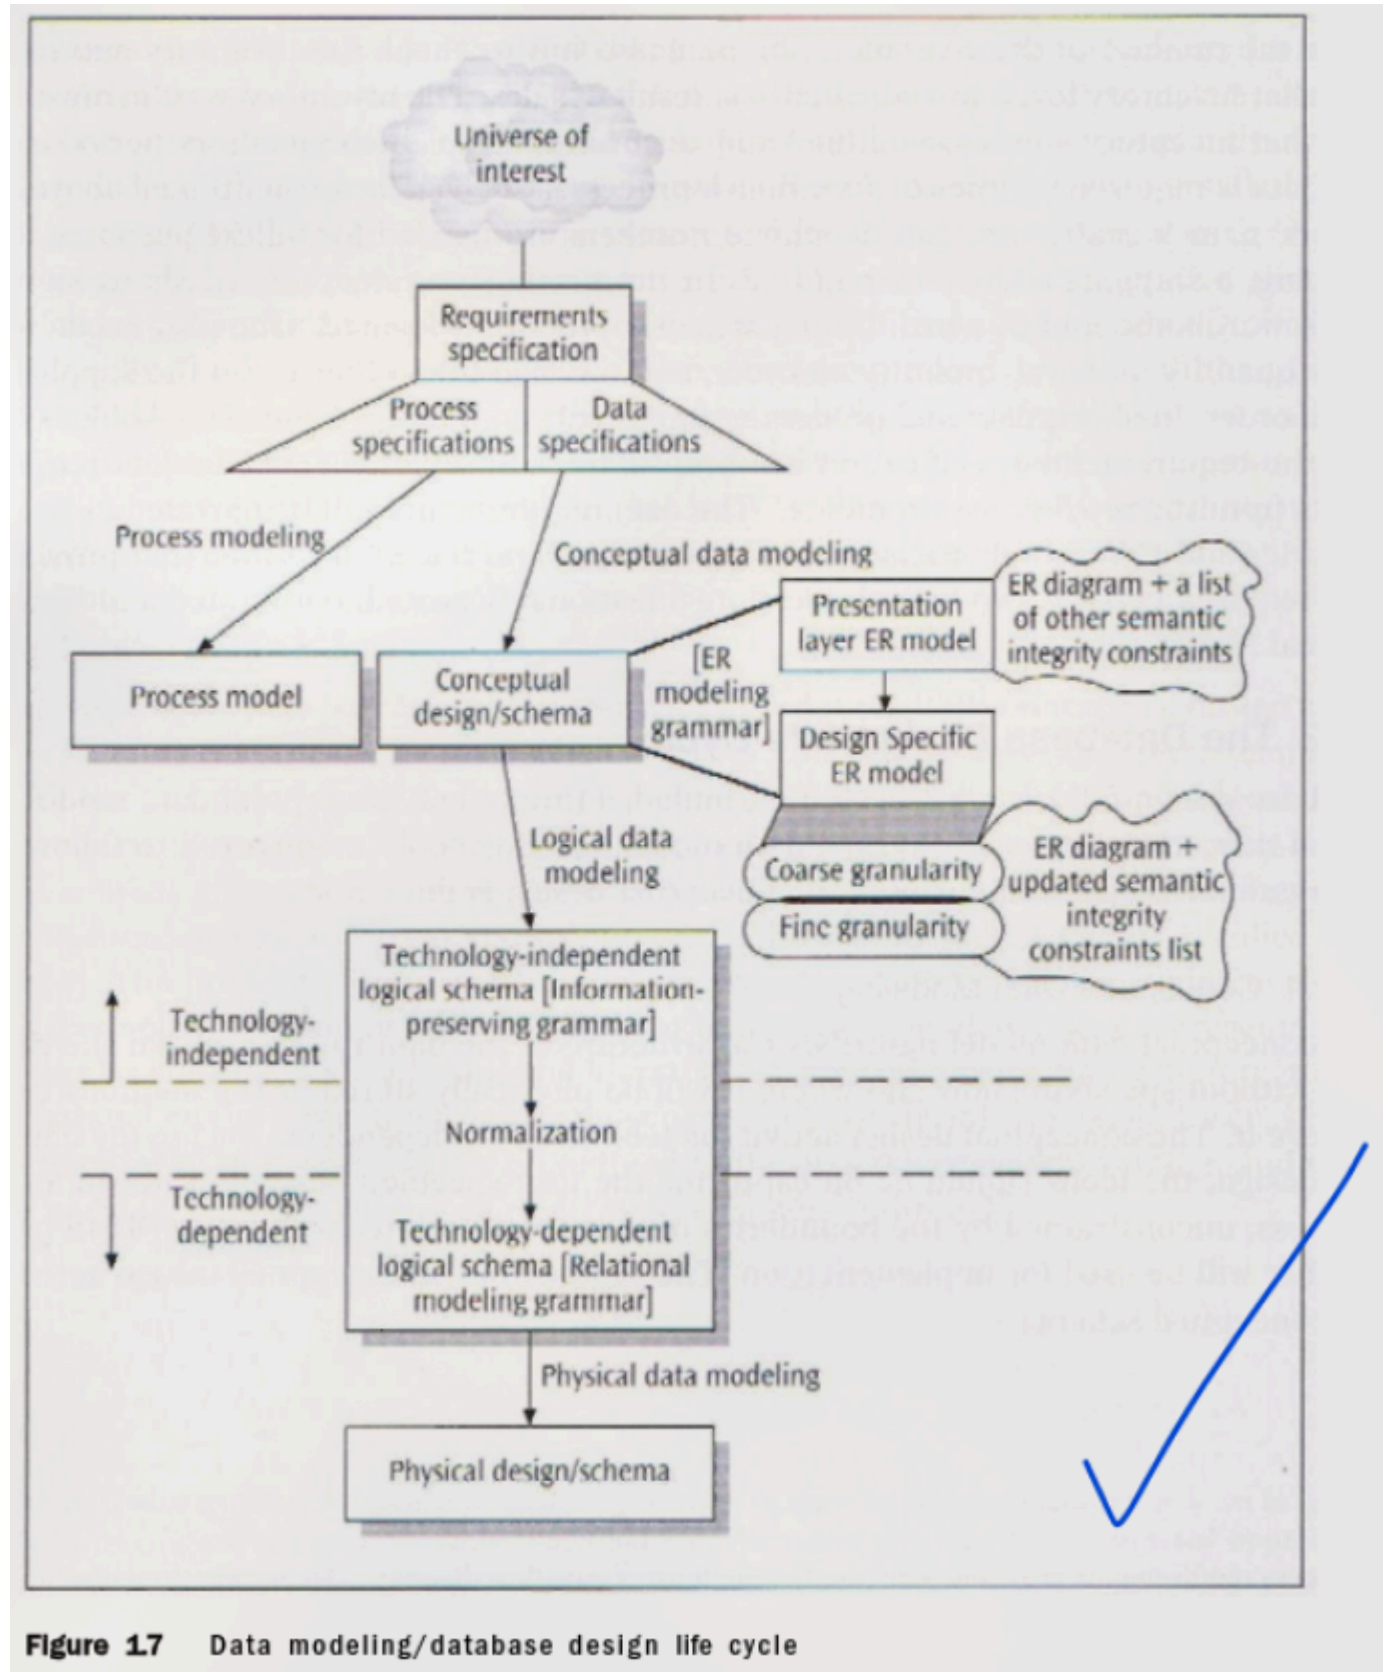
- **Logical Data Independence:** Modify conceptual schema without changing external schemas/applications (e.g., add an attribute).
- **Physical Data Independence:** Modify internal schema without changing conceptual/external schemas (e.g., change file organization, add index).

*(Based on: DBMD UNIT - 1 and 2 Notes, p.2, Figure 1.2 and related text)*

---

### b) Describe the database design life cycle. (7 Marks)

Answer:



**Figure 1.7** Data modeling/database design life cycle

The database design life cycle (DDLC) is a systematic process for designing, implementing, and maintaining a database:

### 1. Requirements Specification (and Planning):

- **Objective:** Understand and document user data requirements and application needs.
- **Activities:** Interview users, review documents, identify objectives, entities, processes, business rules. Define scope.



- **Output:** Detailed requirements specification document.

## 2. Conceptual Data Modeling (Technology-Independent):

- **Objective:** Create a high-level data structure description (entities, attributes, relationships) without physical storage or DBMS specifics.
- **Activities:** Translate requirements to a conceptual model (e.g., ERD, EERD). Identify types, attributes, constraints.
- **Output:** Conceptual schema (e.g., ER/EER diagram), validated with users.

## 3. Logical Data Modeling (Technology-Dependent):

- **Objective:** Transform conceptual schema into a logical model for a chosen data model (typically relational).
- **Activities:** Map ER/EER to tables, attributes to columns, define PKs/FKs. Apply normalization (1NF, 2NF, 3NF, BCNF).
- **Output:** Logical schema (set of normalized relational tables).

## 4. Physical Data Modeling (DBMS-Specific):

- **Objective:** Specify internal storage structures, access paths, file organizations for the chosen DBMS.
- **Activities:** Define column data types, design indexes, consider file organization, clustering, partitioning. Estimate storage, plan security/recovery.
- **Output:** Physical schema (internal schema, DDL statements).

## 5. Implementation and Testing:

- **Objective:** Create and test the database.
- **Activities:** Write DDL, populate data, develop/test applications, perform unit, integration, performance tests.

## 6. Deployment, Operation, and Maintenance:

- **Objective:** Deploy for operational use and maintain over time.
- **Activities:** Deploy DB/apps, monitor performance, tune, regular backups, manage security/access, make modifications.

*(Based on: DBMD UNIT - 1 and 2 Notes, p.5, Figure 1.7 and related text)*

---

### Question 3:

**a) Explain the data models available for the database modelling system. (8 Marks)**

#### **Answer:**

Data models are tools to describe database structure, data types, relationships, and constraints:

#### **1. Hierarchical Data Model:**

- **Structure:** Tree-like; parent-child relationships (1 parent per child).
- **Pros:** Simple for some data, efficient hierarchical access.
- **Cons:** Inflexible (M:N complex), restricted data access.
- **Relevance:** Historical (e.g., IBM's IMS).

## 2. Network Data Model:

- **Structure:** Graph-like; records can have multiple parents/children.
- **Pros:** More flexible than hierarchical, represents M:N better.
- **Cons:** Complex design/management, navigation via pointers.
- **Relevance:** Historical (e.g., IDMS).

## 3. Relational Data Model:

- **Structure:** Data in tables (relations) with rows (tuples) and columns (attributes). Relationships via common columns (PKs/FKs).
- **Pros:** Simple, high data independence, flexible querying (SQL), strong foundation.
- **Cons:** Can be slower for complex joins if unoptimized.
- **Relevance:** Most widely used (e.g., Oracle, MySQL, SQL Server).

## 4. Entity-Relationship (ER) Data Model (Conceptual):

- **Structure:** High-level; real world as entities (objects) and relationships.
- **Representation:** Entities (rectangles), attributes (ovals), relationships (diamonds).
- **Pros:** Easy to understand/communicate design.
- **Cons:** Not directly implemented; a design tool mapped to logical models.
- **Relevance:** Standard for conceptual design.

## 5. Enhanced Entity-Relationship (EER) Data Model (Conceptual):

- **Structure:** Extends ER with superclasses/subclasses, specialization/generalization, inheritance, categorization.
- **Pros:** More precise modeling for complex data.
- **Relevance:** Designing complex databases with hierarchies.

## 6. Object-Oriented Data Model:

- **Structure:** Data as objects (instances of classes) encapsulating data and behavior (methods). Supports inheritance, polymorphism.
- **Pros:** Represents complex types/relationships from OOP.
- **Cons:** Less mature than relational, complex querying.
- **Relevance:** OODBMS, ORDBMS (e.g., CAD/CAM).

## 7. Object-Relational Data Model:

- **Structure:** Hybrid; combines relational features with object-oriented concepts (user-defined types, complex objects).
- **Pros:** Extends relational power, SQL compatibility.
- **Relevance:** Many modern RDBMSs (e.g., PostgreSQL, Oracle).

## 8. NoSQL Data Models (e.g., Document, Key-Value, Column-family, Graph):

- **Structure:** Non-relational; for scalability, flexibility, specific workloads.
  - **Document:** Data in documents (e.g., JSON).
  - **Key-Value:** (key, value) pairs.
  - **Column-family:** Data in columns (good for sparse data).
  - **Graph:** Nodes, edges, properties (for connected data).
- **Pros:** Scalable, flexible schema, good for unstructured/semi-structured data.
- **Cons:** Often eventual consistency, varied query capabilities.
- **Relevance:** Growing for big data, real-time apps.

*(Based on: General knowledge of data models; ER and EER are covered in DBMD UNIT - 1 and 2 Notes, p.6, p.27. Relational model is the basis for much of the notes.)*

---

## b) Illustrate the design issues in ER & EER modelling. (7 Marks)

### Answer:

Effective ER/EER modeling requires addressing several design issues:

### 1. Choosing Entity vs. Attribute:

- **Issue:** Model a concept as an entity or an attribute.
- **Guideline:** Entity if it has own attributes or independent relationships; attribute if a simple property.
- **Example:** "Address" for `Employee`. Attribute if simple; entity if multiple/complex addresses shared.

### 2. Choosing Entity vs. Relationship:

- **Issue:** Represent a concept as an entity or a relationship (esp. M:N with attributes).
- **Guideline:** If an association has descriptive attributes, model as an (associative) entity.
- **Example:** `WORKS_ON` (Employee-Project) with `HoursWorked`. Can be M:N relationship with attribute, or `Assignment` associative entity.

### 3. Binary vs. Higher-Degree Relationships (Ternary, N-ary):

- **Issue:** Degree of a relationship.
- **Guideline:** Use higher-degree sparingly, only if truly irreducible. Often decomposable into binaries.

- **Example:** `SUPPLIES` (Supplier, Part, Project). Ternary if supplier supplies a part *only* for a specific project.
- (See DBMD UNIT - 1 and 2 Notes, p.33-34).

#### 4. Use of Weak Entities:

- **Issue:** When to use a weak entity (existence-dependent, identified via owner).
- **Guideline:** Entity's existence depends on another, PK includes owner's PK.
- **Example:** `Dependents` of an `Employee`. Identified by `(EmployeeID, DependentName)`.
- (See DBMD UNIT - 1 and 2 Notes, p.18).

#### 5. When to Use EER Constructs (Specialization/Generalization, Categorization):

- **Issue:** Use superclass/subclass or categories.
- **Specialization/Generalization (Is-A):** Entity has distinct subgroups (subclasses) with common/specific attributes/relationships. E.g., `Employee` -> `SalariedEmployee`, `HourlyEmployee`.
- **Categorization (Union Type):** Subclass is subset of *union* of distinct entity types. E.g., `VehicleOwner` is `Person` OR `Company`.
- (See DBMD UNIT - 1 and 2 Notes, p.27, p.29, p.32-33).
- **Constraints for Specialization:** Disjointness (d/o), completeness (total/partial).

#### 6. Handling M:N Relationships and Multi-valued Attributes (Design-Specific ER):

- **Issue:** Mapping M:N and multi-valued attributes to relational models.
- **Guideline:** M:N -> associative entity + two 1:N. Multi-valued attribute -> separate entity + 1:N.
- (See DBMD UNIT - 1 and 2 Notes, p.24, p.26).

#### 7. Validation of Conceptual Design (Connection Traps):

- **Issue:** ER structure leading to misinterpretation.
- **Fan Trap:** Ambiguous pathway between entity occurrences due to multiple 1:N relationships fanning out. E.g., `Division` has many `Staff` and many `Branches` (which staff at which branch?).
- **Chasm Trap:** Suggested relationship pathway doesn't exist for all occurrences due to optional participation. E.g., `Branch` -> `Staff` (optional manages) -> `Property`.
- (See DBMD UNIT - 1 and 2 Notes, p.37-40).

**Thoughtful handling of these issues leads to well-structured, accurate database models.**

## UNIT-II

(Select one question from Q4 or Q5)

### Question 4:

**a) Explain the ER model and EER Model to map with logical schema. (8 Marks)**

**Answer:**

Mapping ER/EER models to a logical (typically relational) schema converts conceptual constructs into tables, columns, and keys.

**Mapping ER Model Constructs:**

1. **Strong Entity Types:** Create a table; entity attributes become columns; choose a PK. Composite attributes map to simple component columns.
2. **Weak Entity Types:** Create a table; include attributes. PK = PK of owner entity (as FK) + partial key of weak entity.
3. **1:1 Relationship Types:**
  - **Foreign Key:** Add PK of one entity as a unique FK in the other (prefer side with total participation).
  - **Merged Relation:** If both total participation, consider merging into one table (less common).
  - **Relationship Relation:** Separate table for relationship (PK of one entity as PK, PK of other as unique FK); relationship attributes here.
4. **1:N Relationship Types:** In N-side table, include PK of 1-side entity as FK. Relationship attributes also go in N-side table.
5. **M:N Relationship Types:** Create a new junction/associative table. PK = composite of PKs from both participating entities (as FKs). Relationship attributes go in this new table.
6. **Multi-valued Attributes:** Create a new table. Include PK of original entity (as FK) and the multi-valued attribute. PK of new table = (FK, multi-valued attribute).

**Mapping EER Model Constructs:**

1. **Specialization/Generalization (Superclass/Subclass - SC/sc):**
  - **Option 1 (Multiple Relations - SC and sc's):** Table for SC, table for each sc. SC's PK is PK for each sc (also FK to SC). Good for disjoint subclasses with specific attributes.
  - **Option 2 (Single Relation - SC only):** One table for SC. Include all attributes of SC and all sc's. Use nulls and a "type" attribute. Good for overlapping or few specific sc attributes.
  - **Option 3 (Multiple Relations - sc's only):** Table for each sc, including inherited SC attributes. Only if SC participation is total and disjoint (can cause redundancy).
2. **Specialization Hierarchy/Lattice:** Apply chosen SC/sc mapping recursively. Shared subclass (lattice) might have multiple FKs (if Option 1 used).
3. **Categorization (Union Type):** Create table for category (subclass) with surrogate PK. This PK becomes FK in superclass tables. Complex, often avoided.

*(Based on: DBMD UNIT - 1 and 2 Notes, p.43-44 for ER, p.45, p.47-49 for EER)*


---



b) Construct an ER diagram for the Library Management System covering all major activities and map its logical schema. (7 Marks)

Answer:

ER Diagram for Library Management System:



Interpretation for Logical Schema Mapping:

- PKs: User\_id (Reports), Staff\_id, LoginId, Publisher\_id, User\_ID (Readers), ISBN (Books) are PKs.
- "Reports" entity: Represents transactions (issue/return). Reg\_no likely transaction ID.
- "reserve/return date" M:N relationship (Readers-Books) with attributes -> separate table.
- "maintain" relationship: Staff involved in loan process.
- "Phone no" (Readers): Multi-valued attribute -> separate table.
- "name" (Staff): Split into StaffFirstName, StaffLastName.
- AuthNo (Books): Simple attribute (ideally FK to Authors table not shown).
- YearOfPublication: Attribute of Publisher (e.g., founding year).

## Logical Schema (Relational Tables):

### 1. AuthenticationSystem\_Table

- LoginID (PK, VARCHAR(50))
- Password (VARCHAR(255), Hashed)

### 2. Staff\_Table

- StaffID (PK, VARCHAR(20))
- StaffFirstName (VARCHAR(50), NOT NULL)
- StaffLastName (VARCHAR(50), NOT NULL)
- LoginID (FK, VARCHAR(50), References AuthenticationSystem\_Table(LoginID))

### 3. Publisher\_Table

- PublisherID (PK, VARCHAR(20))
- PublisherName (VARCHAR(100), NOT NULL, UNIQUE)
- YearOfPublication (INT)

### 4. Books\_Table

- ISBN (PK, VARCHAR(13))
- Title (VARCHAR(255), NOT NULL)
- AuthNo (VARCHAR(50))
- Edition (VARCHAR(50))
- Category (VARCHAR(50))
- Price (DECIMAL(10,2))
- PublisherID (FK, VARCHAR(20), References Publisher\_Table(PublisherID))

### 5. Readers\_Table

- UserID (PK, VARCHAR(20))
- FirstName (VARCHAR(50), NOT NULL)
- LastName (VARCHAR(50), NOT NULL)
- Email (VARCHAR(100), UNIQUE)
- Address (VARCHAR(255))

### 6. Reader\_Phones\_Table (For multi-valued Phone no)

- UserID (FK, PK\_part, VARCHAR(20), References Readers\_Table(UserID))
- PhoneNumber (PK\_part, VARCHAR(20))
- Primary Key: (UserID, PhoneNumber)

### 7. Book\_Loan\_Transaction\_Table (From "Reports", "reserve/return date", "Manages", "maintain")

- **TransactionID** (PK, INT or VARCHAR(30)) /\* e.g., **Reg\_no** \*/
- **Reader\_UserID** (FK, VARCHAR(20), References Readers\_Table(UserID), NOT NULL)
- **Book\_ISBN** (FK, VARCHAR(13), References Books\_Table(ISBN), NOT NULL)
- **Issuing\_StaffID** (FK, VARCHAR(20), References Staff\_Table(StaffID), NOT NULL)
- **ReserveDate** (DATE, Nullable)
- **IssueDate** (DATE)
- **DueDate** (DATE, NOT NULL if issued)
- **ReturnDate** (DATE, Nullable)
- **Status** (VARCHAR(20), e.g., 'Issued', 'Returned') /\* **Issue/Return** attribute \*/

#### 8. **Staff\_Reader\_Monitoring\_Table** (From "Keeps track of" M:N relationship)

- **StaffID** (FK, PK\_part, VARCHAR(20), References Staff\_Table(StaffID))
- **Reader\_UserID** (FK, PK\_part, VARCHAR(20), References Readers\_Table(UserID))
- **MonitoringStartDate** (DATE, Optional)
- **Primary Key:** (**StaffID**, **Reader\_UserID**)

*(Based on ER design principles and mapping rules from DBMD UNIT - 1 and 2 Notes)*

---

### Question 5:

#### a) Define Normalization. Explain the types of Normalization. (8 Marks)

##### Answer:

##### Definition of Normalization:

Normalization is organizing database data to reduce redundancy and improve data integrity. It involves decomposing tables into smaller, well-structured ones with defined relationships.

**Goals:** Minimize redundancy, eliminate Insert/Update/Delete anomalies, ensure logical data dependencies, produce flexible/maintainable design.

##### Types of Normalization (Normal Forms):

##### 1. First Normal Form (1NF):

- **Rule:** All attribute values are atomic (single, indivisible). No repeating groups or multi-valued attributes in a cell.
- *(See DBMD UNIT - 1 and 2 Notes, p.51-52).*

##### 2. Second Normal Form (2NF):

- **Rule:** In 1NF, and every non-primary-key attribute is fully functionally dependent on the *entire* primary key (no partial dependencies).
- *(See DBMD UNIT - 1 and 2 Notes, p.52-53).*

### 3. Third Normal Form (3NF):

- **Rule:** In 2NF, and no transitive dependencies (non-key attribute depends on another non-key attribute).
- (See DBMD UNIT - 1 and 2 Notes, p.53-54).

### 4. Boyce-Codd Normal Form (BCNF):

- **Rule:** In 3NF, and for every non-trivial functional dependency  $X \rightarrow Y$ ,  $X$  must be a superkey. Stricter than 3NF.
- (See DBMD UNIT - 1 and 2 Notes, p.54-55).

### 5. Fourth Normal Form (4NF):

- **Rule:** In BCNF, and no non-trivial multi-valued dependencies (MVDs  $X \twoheadrightarrow Y$ ), unless  $X$  is a superkey. Separates independent multi-valued facts.
- (See DBMD UNIT - 1 and 2 Notes, p.55-57).

### 6. Fifth Normal Form (5NF / Project-Join Normal Form - PJ/NF):

- **Rule:** In 4NF, and no join dependencies (JDs) not implied by candidate keys. Ensures no further lossless decomposition possible (except by candidate keys).
- (See DBMD UNIT - 1 and 2 Notes, p.57-58).

3NF or BCNF is often sufficient for most designs.

---

## b) Explain the Mapping of higher degree relationships. (7 Marks)

### Answer:

Higher-degree relationships (ternary, n-ary) involve three or more entity types in a single relationship instance.

### General Approach: Associative Entity (Junction Table)

Map any higher-degree relationship by creating a new relation (associative entity/junction table) for the relationship.

### Steps for Mapping an N-ary Relationship:

1. **Create New Relation:** For n-ary relationship  $R$ , create table  $S$ .
2. **Include Foreign Keys:** For each participating entity  $E_1, \dots, E_n$ , include  $PK(E_i)$  as an FK in  $S$ .
3. **Primary Key of  $S$ :** Typically, the combination of all these FKs:  $\{PK(E_1), \dots, PK(E_n)\}$ .
4. **Map Relationship Attributes:** Attributes of  $R$  become columns in  $S$ .
5. **Cardinality Constraints:** Exact (min,max) constraints for n-ary relationships are hard to enforce with only PK/FK; may need application logic or complex checks.

**Example: Ternary Relationship** `SCHEDULE` (Instructor, Course, Quarter)

Entities: `INSTRUCTOR(InstructorID_PK, ...)`, `COURSE(CourseID_PK, ...)`, `QUARTER(QuarterID_PK, ...)`

Relationship `SCHEDULE` has attribute `RoomNo`.

**Mapping** `SCHEDULE`:

```
CREATE TABLE SCHEDULE_TABLE (  
    InstructorID_FK INT,          -- FK to INSTRUCTOR  
    CourseID_FK VARCHAR(10),     -- FK to COURSE  
    QuarterID_FK VARCHAR(10),    -- FK to QUARTER  
    RoomNo VARCHAR(10),          -- Attribute of SCHEDULE  
    PRIMARY KEY (InstructorID_FK, CourseID_FK, QuarterID_FK),  
    FOREIGN KEY (InstructorID_FK) REFERENCES INSTRUCTOR(InstructorID_PK),  
    FOREIGN KEY (CourseID_FK) REFERENCES COURSE(CourseID_PK),  
    FOREIGN KEY (QuarterID_FK) REFERENCES QUARTER(QuarterID_PK)  
);
```

**Considerations:**

- **Decomposition:** Evaluate if the higher-degree relationship can be losslessly decomposed into binary ones. Preferable if semantics are preserved.
- **Conceptual Model First:** (DBMD Notes, p.50) Suggests decomposing into a gerund (associative) entity in the conceptual model first, then mapping this entity and its binary relationships.
- **Aggregation:** Distinct concept where a relationship is treated as a higher-level entity.

The key is to accurately capture all instances and attributes of the higher-degree relationship.

---

## UNIT-III

(Select one question from Q6 or Q7)

**Question 6:**

**a) Describe the database creation using SQL with help of suitable examples. (8 Marks)**

**Answer:**

Database creation in SQL mainly uses the `CREATE TABLE` DDL statement to define a new table, its columns, data types, and constraints.

**Syntax of** `CREATE TABLE` (Simplified):

```
CREATE TABLE table_name (  
    column_name_1 data_type [column_constraints],  
    column_name_2 data_type [column_constraints],
```



```
...  
[table_constraints]  
);
```

### Key Components:

- `table_name`: Name of the new table.
- `column_name`: Name of a column.
- `data_type`: Type of data (e.g., `INTEGER`, `VARCHAR(n)`, `DATE`, `DECIMAL(p,s)`).
- `column_constraints`: Apply to individual columns (e.g., `NOT NULL`, `UNIQUE`, `PRIMARY KEY`, `CHECK (condition)`, `DEFAULT default_value`, `REFERENCES referenced_table`).
- `table_constraints`: Apply to one or more columns (e.g., `PRIMARY KEY (cols)`, `UNIQUE (cols)`, `FOREIGN KEY (cols) REFERENCES ...`, `CHECK (condition)`).

### Example: Creating Tables for a Medical System

(Based on Figure 10.1a, Box 1, Box 2, Box 3 in notes)

#### 1. Patient Table:

```
CREATE TABLE Patient (  
    Pat_p_alpha CHAR(2) NOT NULL,  
    Pat_p_num CHAR(5) NOT NULL,  
    Pat_name VARCHAR(100) NOT NULL,  
    Pat_gender CHAR(1) CHECK (Pat_gender IN ('M', 'F')),  
    Pat_age SMALLINT CHECK (Pat_age >= 0 AND Pat_age <= 120),  
    Pat_admit_date DATE,  
    Pat_wing CHAR(1),  
    Pat_room_num INT,  
    Pat_bed CHAR(1) CHECK (Pat_bed IN ('A', 'B')),  
    PRIMARY KEY (Pat_p_alpha, Pat_p_num)  
);
```

#### 2. Medication Table:

```
CREATE TABLE Medication (  
    Med_code CHAR(5) PRIMARY KEY,  
    Med_name VARCHAR(100) NOT NULL UNIQUE,  
    Med_unit_price DECIMAL(5,2) CHECK (Med_unit_price > 0),  
    Med_qty_onhand INT DEFAULT 0,  
    Med_qty_onorder INT DEFAULT 0,  
    CONSTRAINT chk_med_stock CHECK ((Med_qty_onhand + Med_qty_onorder) BETWEEN 0  
AND 5000)  
);
```

#### 3. Order Table (linking Patient and Medication):

```
CREATE TABLE Orders (
    Ord_rx_num CHAR(13) PRIMARY KEY,
    Ord_pat_p_alpha CHAR(2) NOT NULL,
    Ord_pat_p_num CHAR(5) NOT NULL,
    Ord_med_code CHAR(5) NOT NULL,
    Ord_dosage SMALLINT DEFAULT 1 CHECK (Ord_dosage IN (1, 2, 3)),
    Ord_freq SMALLINT DEFAULT 1 CHECK (Ord_freq IN (1, 2, 3)),
    FOREIGN KEY (Ord_pat_p_alpha, Ord_pat_p_num) REFERENCES Patient(Pat_p_alpha,
Pat_p_num)
        ON DELETE CASCADE ON UPDATE CASCADE,
    FOREIGN KEY (Ord_med_code) REFERENCES Medication(Med_code)
        ON DELETE RESTRICT ON UPDATE RESTRICT
);
```

This shows tables with various data types, PKs (simple/composite), FKs with referential actions, CHECK, and DEFAULT constraints.

*(Based on: DBMD UNIT - 3 and 4 Notes, p.1-8)*

---

## b) Explain briefly DDL and DML with syntax and suitable examples. (7 Marks)

**Answer:**

### DDL (Data Definition Language):

Defines, modifies, and removes database structures (schemas) and objects.

- **Key DDL Commands:**

1. **CREATE**: Creates new DB objects (tables, views).

- **Syntax (Table):** `CREATE TABLE table_name (col1 datatype, ...);`

- **Example:** `CREATE TABLE Products (ProdID INT PRIMARY KEY, Name VARCHAR(50));`

2. **ALTER**: Modifies existing DB objects.

- **Syntax (Add Column):** `ALTER TABLE table_name ADD new_column datatype;`

- **Example:** `ALTER TABLE Products ADD Price DECIMAL(10,2);`

3. **DROP**: Deletes existing DB objects.

- **Syntax (Table):** `DROP TABLE table_name;`

- **Example:** `DROP TABLE Products;`

4. **TRUNCATE**: Removes all rows from a table (structure remains); faster than **DELETE** for all rows.

- **Syntax (Table):** `TRUNCATE TABLE table_name;`

- **Example:** `TRUNCATE TABLE TempData;`

## DML (Data Manipulation Language):

Retrieves, inserts, updates, and deletes data within tables.

- **Key DML Commands:**

1. **SELECT**: Retrieves data.

- **Syntax:** `SELECT col_list FROM table_name [WHERE condition];`

- **Example:** `SELECT Name, Price FROM Products WHERE Price > 100;`

2. **INSERT**: Adds new rows.

- **Syntax:** `INSERT INTO table_name (col1, col2) VALUES (val1, val2);`

- **Example:** `INSERT INTO Products (ProdID, Name, Price) VALUES (1, 'Laptop', 1200);`

3. **UPDATE**: Modifies existing data.

- **Syntax:** `UPDATE table_name SET col1 = val1 [WHERE condition];`

- **Example:** `UPDATE Products SET Price = 1150 WHERE ProdID = 1;`

4. **DELETE**: Removes rows.

- **Syntax:** `DELETE FROM table_name [WHERE condition];`

- **Example:** `DELETE FROM Products WHERE ProdID = 1;`

(Based on: DBMD UNIT - 1 and 2 Notes, p.4; DBMD UNIT - 3 and 4 Notes, p.1, p.8, p.12)

---

### Question 7:

**a) Describe the cursor and type of cursor. What is the need of cursor in database programming?**

(8 Marks)

**Answer:**

**Cursor:**

A cursor is a database control structure enabling row-by-row traversal over a result set returned by an SQL query. It acts as a pointer to a specific row, facilitating procedural processing of set-based SQL results.

**Need for Cursors:** (DBMD UNIT - 3 and 4 Notes, p.18)

SQL is set-oriented, while host languages (C, Java) are record-oriented, creating an **impedance mismatch**. Cursors bridge this by allowing applications to:

1. Define a row set (from an SQL query).
2. Iterate through this set, fetching one row at a time into host variables.
3. Potentially update/delete the current row (with updatable cursors).

**Operations on Cursors:** (DBMD UNIT - 3 and 4 Notes, p.18)

**DECLARE** (defines cursor with query), **OPEN** (executes query, positions cursor), **FETCH** (retrieves row,

advances cursor), `UPDATE ... WHERE CURRENT OF` (modifies current row), `DELETE ... WHERE CURRENT OF` (deletes current row), `CLOSE` (releases resources).

### Types of Cursors: (DBMD UNIT - 3 and 4 Notes, p.19-20)

1. **Read-Only Cursor:** Allows only fetching data. (Default or `FOR READ ONLY`).
2. **Updatable Cursor:** Allows fetching and modifying/deleting current row (`FOR UPDATE [OF column_list]`). Query usually simple.
3. **Forward-Only Cursor (Non-Scrollable):** Default; rows fetched sequentially (first to last). Most efficient.
4. **Scrollable Cursor:** Flexible movement (`NEXT`, `PRIOR`, `FIRST`, `LAST`, `ABSOLUTE n`, `RELATIVE n`). (`SCROLL CURSOR`).
5. **Insensitive Cursor (Snapshot Cursor):** Operates on a temporary copy of data at open time. Changes by others not visible. (`INSENSITIVE CURSOR`). Provides static view.
6. **Sensitive Cursor:** Attempts to reflect changes by others after open. Behavior varies, can be complex.
7. **Keyset-Driven Cursor:** A type of sensitive cursor. Keys of qualifying rows fixed at open. Non-key value changes visible; deleted rows appear as "holes"; new qualifying rows usually not seen.

(Note: Static/Dynamic cursors map broadly to Insensitive/highly Sensitive cursors.)

---

### b) What is database trigger? Explain the types of Trigger. (7 Marks)

**Answer:**

#### Database Trigger:

A procedural code (SQL block/stored procedure) automatically executed by the DBMS in response to specific DML events (e.g., `INSERT`, `UPDATE`, `DELETE`) or DDL events (`CREATE`) on a table/view.

#### ECA Model (Event-Condition-Action): (DBMD UNIT - 3 and 4 Notes, p.21)

- **Event:** The DML/DDDL operation causing trigger to fire.
- **Condition (Optional):** Boolean expression; action executes if true.
- **Action:** Code executed when event occurs and condition met.

#### Syntax (Simplified Generic): (DBMD UNIT - 3 and 4 Notes, p.21)

```
CREATE TRIGGER trigger_name {BEFORE | AFTER} {INSERT | DELETE | UPDATE [OF cols]}  
ON table_name [FOR EACH ROW] [WHEN (condition)] BEGIN ... END;
```

#### Types of Triggers:

1. **Row-Level vs. Statement-Level Trigger:**

- **Row-Level (FOR EACH ROW):** Fires once per affected row. Can access :OLD/:NEW row values.  
Ex: Auditing each updated row.
- **Statement-Level (Default):** Fires once per DML statement. Cannot access :OLD/:NEW. Ex:  
Alert on any delete attempt from critical table.

## 2. Timing (BEFORE vs. AFTER):

- **BEFORE Trigger:** Action executes *before* triggering DML. Uses: Validate/modify new data, prevent operations.
- **AFTER Trigger:** Action executes *after* triggering DML and constraints. Uses: Auditing, complex integrity, propagating changes.

## 3. INSTEAD OF Triggers: (DBMD UNIT - 3 and 4 Notes, p.21)

- For views (especially non-updatable ones). Fires *instead* of DML on view. Defines how to apply operation to base tables.

## 4. DDL Triggers: (DBMD UNIT - 3 and 4 Notes, p.22)

- Fire on DDL events (CREATE TABLE). Uses: Audit schema changes, enforce naming. (DBMS-specific support).

## 5. Logon/Logoff Triggers (System-Level): (DBMD UNIT - 3 and 4 Notes, p.22)

- Fire on user connect/disconnect. Uses: Audit sessions, set parameters. (DBMS-specific).

## 6. Database Event Triggers (Server-Level): (DBMD UNIT - 3 and 4 Notes, p.22)

- Fire on DB events (startup, shutdown, errors). Uses: Admin tasks, custom logging. (DBMS-specific).

## 7. Compound Triggers (Oracle specific): (DBMD UNIT - 3 and 4 Notes, p.22)

- Defines actions for multiple timing points (BEFORE STATEMENT, BEFORE EACH ROW, etc.) for one DML event in a single trigger.

Triggers are powerful but can add complexity; use carefully.

## UNIT-IV

(Select one question from Q8 or Q9)

### Question 8:

a) Explain the Indexing and its methods with the help of suitable examples. (8 Marks)

**Answer:**

**Indexing:**

A database technique to speed up row retrieval. An index is a separate data structure (e.g., B+-tree) storing copies of one or more columns (index key) with pointers (rowids) to actual data rows. It helps locate relevant rows quickly without full table scans.



**Why Indexing?** Avoids inefficient full table scans for large tables by providing a shortcut.

### Methods/Types of Indexing:

1. **Primary Index (often Clustered):** Index key specifies data file's sequential order. Table physically ordered by this key. Max one per table.
  - **Example:** Index on `EmployeeID` (PK), employee records physically stored in `EmployeeID` order.
2. **Clustered Index:** Physical order of table rows matches index key order.
  - **Example:** `CREATE CLUSTERED INDEX IX_Orders_Date ON Orders(OrderDate);` `Orders` table rows physically sorted by `OrderDate`. Efficient for range queries on `OrderDate`.
3. **Secondary Index (Non-Clustered Index):** Index key order differs from physical data order. Index entries point to data rows. Multiple per table.
  - **Example:** `CREATE INDEX IX_Emp_LName ON Employees(LastName);` Allows quick lookups by `LastName`.
4. **B+-Tree Index:** Most common; balanced tree. Leaf nodes contain (key, pointer) pairs, linked sequentially. Efficient for equality and range searches.
  - **Example:** Default index type, e.g., `CREATE INDEX IX_Prod_Price ON Products(Price);`
5. **Hash Index:** Uses hash function to find bucket/page for an index entry.
  - **Example:** Index on `CustomerEmail`. Fast for `WHERE CustomerEmail = '...';` Not for range queries.
6. **Unique Index:** Enforces no duplicate values for indexed column(s). PK indexes are unique.
  - **Example:** `CREATE UNIQUE INDEX UQ_Emp_Email ON Employees(Email);`
7. **Composite Index (Multi-column Index):** Index on two or more columns. Column order is key.
  - **Example:** `CREATE INDEX IX_Orders_CustDate ON Orders(CustomerID, OrderDate);` Useful for filters on `CustomerID` or `CustomerID AND OrderDate`.
8. **Covering Index:** Non-clustered index containing all columns for a query (in `SELECT` & `WHERE`). Query answered from index (index-only scan).
  - **Example:** For `SELECT OrderDate, Amount FROM Orders WHERE CustID = 1;`, index on `(CustID, OrderDate, Amount)` is covering.

(Based on: DBMD UNIT - 3 and 4 Notes, p.23-27)

---

### b) Discuss about the database security used in the database modelling. (7 Marks)

#### Answer:

Integrating security into database modeling is crucial for building secure systems.

### Security Considerations During Modeling Phases:

## 1. Conceptual Level (ER/EER & Requirements):

- **Identify Sensitive Data:** Pinpoint entities/attributes needing protection (e.g., PII, financial data).
- **Define Access Privileges Conceptually (User Roles):** Document roles (HR, Sales) and their intended data access (read, create, modify, delete).
- **Consider Views for Security:** Note needs for subset views (hiding columns, showing aggregated data) for specific user groups. E.g., view for average salaries, not individual.
- **Data Ownership:** Clarify who owns data, informing access granting authority.

## 2. Logical Level (Relational Schema Design):

- **View Design:** Implement views for:
  - **Column-level security:** Hide sensitive columns.
  - **Row-level security:** Filter rows via `WHERE` clause based on user identity/attributes (e.g., manager sees only their region's orders).
- **Granular Privilege Planning:** Plan SQL privileges (`SELECT`, `INSERT`, etc.) for roles on tables/views, informing DCL.
- **Normalization/Integrity:** Proper normalization helps maintain data integrity (a facet of security).

## 3. Physical Level (DBMS Implementation - Informed by Modeling):

- **GRANT/REVOKE (DAC):** Implement planned access privileges.
- **Role-Based Access Control (RBAC):** Create roles, grant privileges to roles, assign users to roles.
- **Stored Procedures:** Encapsulate sensitive DML; grant `EXECUTE` on procedure, not direct table access.
- **Triggers for Auditing:** Implement triggers to log access to sensitive data.
- **Encryption:** Plan for column-level or Transparent Data Encryption (TDE) for highly sensitive data.

Considering security at each modeling stage ensures controls align with requirements, protect data, and enforce least privilege.

*(Based on: DBMD UNIT - 3 and 4 Notes, p.31-32)*

---

### Question 9:

**Write short notes on (5x3 = 15 Marks)**

#### a) Clustering

#### Answer:

Clustering is physically storing related data records close together on disk (ideally same/adjacent

blocks) to minimize Disk I/O and improve query performance for accessing these records together. Data is typically organized by a **clustering key**.

**Types:** (DBMD UNIT - 3 and 4 Notes, p.28)

1. **Intra-Table Clustering (Clustered Index):** Physical row order within a table matches its clustered index order. A table has at most one. Efficient for range queries on clustering key.
  - **Example:** `Orders` table with clustered index on `OrderDate`; records physically sorted by `OrderDate`.
2. **Inter-Table Clustering (Co-clustering):** Physically interleaving rows from related tables based on a common join key (PK-FK). Speeds up frequent joins.
  - **Example:** Storing `OrderItem` records near their parent `Order` record.

**Benefits:** Reduced Disk I/O, faster joins (inter-table), efficient range queries (intra-table).

**Drawbacks:** Slower DML (order maintenance), only one clustered index/table (intra-table).

(Based on: DBMD UNIT - 3 and 4 Notes, p.23, p.28-29)

---

## b) De-normalization

### Answer:

Denormalization is intentionally introducing controlled redundancy into a relational schema (adding duplicate/grouped data) to improve read performance (query speed) by reducing joins or pre-calculating values. It's the reverse of normalization.

**Rationale:** (DBMD UNIT - 3 and 4 Notes, p.29-30)

Applied selectively when highly normalized schemas lead to slow queries due to many joins, and indexing/tuning are insufficient.

**Techniques:** (DBMD UNIT - 3 and 4 Notes, p.30)

1. **Pre-joining Tables:** Adding attributes from "one-side" to "many-side" table (e.g., `DepartmentName` in `EMPLOYEE` table).
2. **Storing Derived/Calculated Values:** Pre-computing expensive values (e.g., `OrderTotal` in `ORDERS` table).
3. **Combining Tables:** Merging tables with 1:1 or tight, frequently accessed 1:N relationships.
4. **Repeating Groups (Limited):** Using multiple columns for fixed, few similar attributes (e.g., `Phone1`, `Phone2`).
5. **Creating Reporting Tables/Data Marts:** Separate, denormalized tables for analytics.

**Trade-offs:**

- **Benefit:** Faster reads/queries.
- **Cost:** Increased storage, data inconsistency risk (needs careful sync), slower updates, complex DML.

**Guideline:** Apply judiciously *after* normalization, only for clear performance bottlenecks unresolved by other means.

*(Based on: DBMD UNIT - 3 and 4 Notes, p.29-30)*

---

## c) Database Tuning

### **Answer:**

Database tuning is systematically optimizing database system aspects to improve performance and meet user requirements. It's an iterative process to resolve bottlenecks.

**Key Areas:** (DBMD UNIT - 3 and 4 Notes, p.22, p.30-31)

#### 1. Tuning Conceptual Schema (Logical Design):

- Normalization/Denormalization evaluation.
- Vertical/Horizontal Partitioning.

#### 2. Tuning Queries and Views:

- Rewriting inefficient SQL (sargable predicates, simpler logic).
- Optimizing view definitions.
- Analyzing execution plans, ensuring up-to-date statistics.

#### 3. Tuning Physical Design (Indexing and Storage):

- Index selection/management (create appropriate, drop unused, rebuild).
- Clustering decisions.
- File organization/placement, disk space management.

4. **Tuning Application Code:** Optimizing DB interaction (e.g., batch updates, connection management).

5. **Tuning DBMS Parameters:** Adjusting config (memory, I/O, concurrency). DBMS-specific.

6. **Tuning Hardware and OS:** Ensuring sufficient resources, optimizing OS settings.

**Iterative Process:** (DBMD UNIT - 3 and 4 Notes, p.22)

1. **Monitor:** Track performance.
2. **Identify Bottlenecks:** Pinpoint poor performance areas.
3. **Diagnose:** Find root cause.
4. **Implement Changes:** Apply tuning measures.

5. **Measure:** Evaluate impact; keep if improved, else revert/retry.

**Workload Analysis:** (DBMD UNIT - 3 and 4 Notes, p.22)

Critical first step: understand query/update types, frequencies, performance goals, accessed data.

*(Based on: DBMD UNIT - 3 and 4 Notes, p.22, p.29-31)*

---